

Selected Subjects in Robotics and AI HW1

Daniel Haim Breger, 316136944

Technion

(Date: November 20, 2025)

1 Maximum Likelihood Estimation for Non-Gaussian Distribution

Let $(X_1, X_2, \dots, X_n)$ be i.i.d samples from a Rayleigh distribution with unknown scale parameter $(\theta > 0)$:

$$f(x, \theta) = \frac{x}{\theta^2} e^{-x^2/(2\theta^2)}, \quad x \geq 0$$

1. Derive the maximum likelihood estimator $\hat{\theta}_{ML}$.
2. Show that it is a consistent estimator of θ .
3. Determine whether $\hat{\theta}_{ML}$ is biased. If it is, derive the bias expression.
4. Compute the Fisher Information and Cramer-Rao lower bound for estimating θ and check if $\hat{\theta}_{ML}$ attains it asymptotically.

1.1 Deriving MLE

The log-likelihood is

$$l(\theta) = \sum_{i=1}^n \left(\log X_i - 2 \log \theta - \frac{X_i^2}{2\theta^2} \right) \quad (1)$$

The definition of the MLE is

$$\hat{\theta}_{MLE} = \arg \max L(\theta, x) \quad (2)$$

To find the argmax we'll differentiate the log-likelihood:

$$\frac{dl}{d\theta} = \sum_{i=1}^n \left(-\frac{2}{\theta} + \frac{X_i^2}{\theta^3} \right) = \frac{-2n}{\theta} + \frac{1}{\theta^3} \sum_{i=1}^n X_i^2 \quad (3)$$

we require the derivative is equal to zero and find

$$\frac{dl}{d\theta} = \frac{-2n}{\theta} + \frac{1}{\theta^3} \sum_{i=1}^n X_i^2 = 0 \quad (4)$$

$$\rightarrow -2n\theta^2 + \sum_{i=1}^n X_i^2 = 0 \quad (5)$$

$$\rightarrow \theta_{\max} = \sqrt{\frac{1}{2n} \sum_{i=1}^n X_i^2} \quad (6)$$

verifying the point we found is indeed a maximum

$$\left. \frac{d^2l}{d\theta^2} \right|_{\theta_{\max}} = \left. \frac{2n}{\theta^2} - \frac{3}{\theta^4} \sum_{i=1}^n X_i^2 \right|_{\theta_{\max}} \quad (7)$$

$$= \frac{2n}{\theta_{\max}^2} - \frac{3}{\theta_{\max}^4} \cdot 2n\theta_{\max}^2 \quad (8)$$

$$= \frac{2n - 6n}{\theta_{\max}^4} = -\frac{4n}{\theta_{\max}^4} < 0 \quad (9)$$

Therefore it is indeed a maximum, so our MLE is

$$\hat{\theta}_{MLE} = \sqrt{\frac{1}{2n} \sum_{i=1}^n X_i^2} \quad (10)$$

1.2 Consistency Check

An estimator is called consistent if

$$\lim_{n \rightarrow \infty} \hat{\theta}_n = \theta \quad (11)$$

So we'll look at the limit

$$\lim_{n \rightarrow \infty} \hat{\theta}_{MLE} = \lim_{n \rightarrow \infty} \sqrt{\frac{1}{2n} \sum_{i=1}^n X_i^2} \quad (12)$$

Using the law of large numbers we can say

$$\lim_{n \rightarrow \infty} \sqrt{\frac{1}{2n} \sum_{i=1}^n X_i^2} = \sqrt{\frac{1}{2n} \sum_{i=1}^n \mathbb{E}(X_i^2)} = \sqrt{\frac{1}{2n} 2n\theta^2} = \theta \quad (13)$$

Using the 2nd moment of the Rayleigh distribution. Thus it is a consistent estimator.

1.3 Bias Check

The bias is defined as

$$\text{Bias}(\hat{\theta}) = \mathbb{E}(\hat{\theta}) - \theta \quad (14)$$

thus for our MLE estimator

$$\text{Bias}(\hat{\theta}) = \mathbb{E} \left(\sqrt{\frac{1}{2n} \sum_{i=1}^n X_i^2} \right) - \theta \quad (15)$$

$$= \frac{1}{\sqrt{2n}} \mathbb{E} \left(\sqrt{\sum_{i=1}^n X_i^2} \right) - \theta \quad (16)$$

to find the expectation value we'll notice that if $X_i \sim \text{Rayleigh}(\theta)$ then $\sum X_i^2 \sim \Gamma(n, 2\theta^2)$. We can use the formula for the moments of the gamma distribution for the 1/2-th moment

$$\mathbb{E}(\sqrt{\sum X_i^2}) = \sqrt{2\theta} \frac{\Gamma(n + \frac{1}{2})}{\Gamma(n)} \quad (17)$$

and inserting back we find

$$\text{Bias}(\hat{\theta}) = \frac{1}{\sqrt{2n}} \mathbb{E} \left(\sqrt{\sum_{i=1}^n X_i^2} \right) - \theta \quad (18)$$

$$= \frac{1}{\sqrt{2n}} \sqrt{2\theta} \frac{\Gamma(n + \frac{1}{2})}{\Gamma(n)} - \theta \quad (19)$$

$$= \theta \left(\frac{1}{\sqrt{n}} \frac{\Gamma(n + \frac{1}{2})}{\Gamma(n)} - 1 \right) \quad (20)$$

The expression in the brackets is non-zero so the estimator is biased.

1.4 Fisher Information and CRLB

The Fisher Information is defined by

$$I = -\mathbb{E} \left(\frac{\partial^2}{\partial \theta^2} \log f \right) = -\mathbb{E} \left(\frac{\partial^2}{\partial \theta^2} \left(\log x - 2 \log \theta - \frac{x^2}{2\theta^2} \right) \right) \quad (21)$$

$$= -\mathbb{E} \left(\frac{\partial}{\partial \theta} \left(-\frac{2}{\theta} + \frac{x^2}{\theta^3} \right) \right) \quad (22)$$

$$= -\mathbb{E} \left(\frac{2}{\theta^2} - \frac{3x^2}{\theta^4} \right) \quad (23)$$

$$= - \left(\frac{2}{\theta^2} - \frac{3}{\theta^4} \mathbb{E}(x^2) \right) \quad (24)$$

$$= - \left(\frac{2}{\theta^2} - \frac{3}{\theta^4} 2\theta^2 \right) \quad (25)$$

$$= \frac{-2 + 6}{\theta^2} = \frac{4}{\theta^2} \quad (26)$$

And since the Fisher information is based on expectation value it is linear and thus the fisher information of n i.i.d variables is

$$I_n = \frac{4n}{\theta^2} \quad (27)$$

The CRLB is therefore

$$\text{Var}(\hat{\theta}_{MLE}) \geq \frac{\theta^2}{4n} \quad (28)$$

The MLE estimator is indeed asymptotically efficient as we saw in the lecture, but to show this we can use the asymptotic normality property

$$\sqrt{n} \left(\hat{\theta}_{MLE} - \theta \right) \rightarrow N\left(0, \frac{\theta^2}{4n}\right) \quad (29)$$

therefore asymptotically

$$\text{Var}(\hat{\theta}_{MLE}) = \frac{\theta^2}{4n} \quad (30)$$

which is indeed the CRLB.

2 Maximum A Posteriori Estimation with a Specific Prior

Assume the same Rayleigh distributed observations as in question 1, but now suppose the prior knowledge about θ is modeled as an inverse-gamma distribution

$$p(\theta) = \frac{\beta^\alpha}{\Gamma(\alpha)} \theta^{-(\alpha+1)} e^{-\beta/\theta}, \quad \theta > 0$$

1. Derive the MAP estimator $\hat{\theta}_{MAP}$
2. Compare $\hat{\theta}_{ML}$ and $\hat{\theta}_{MAP}$ in terms of bias and asymptotic behavior.
3. Analyze how the choice of prior parameters α and β influences the MAP estimate.
4. Discuss the limiting cases $(\alpha, \beta \rightarrow 0)$ and $(\alpha, \beta \rightarrow \infty)$ and interpret their effects on the posterior estimate.

3 Cramer-Rao Lower Bound and Efficiency

Let $(X_1, X_2, \dots, X_n)$ be i.i.d samples from an exponential distribution

$$f(x, \lambda) = e^{-\lambda x}, \quad x \geq 0$$

where $\lambda > 0$ is unknown.

1. Derive the ML estimator $\hat{\lambda}_{ML}$.
2. Compute the Fisher Information.
3. Derive the CRLB for any unbiased estimator of λ .
4. Show that $\hat{\lambda}_{ML}$ achieves the CRLB and hence is an effective estimator.

4 Neural Network Regression Implementation

The implementation for this is found in the included file `q4.ipynb`. This section will cover the discussion on comparison between the NN estimator and the linear regression.

For personal curiosity reasons we ran multiple NN architectures to see the effects of different depths and layer sizes. Of all the runs which can be found in detail in the jupyter notebook, the better performing ones all were similar in size and depth.

Method	MSE
Linear Regression	0.524
NN (8-32-32-1)	0.270
NN (8-32-16-1)	0.276
NN (8-16-16-16-1)	0.275

Table 1 – Different estimators for mean house value used on the scikit-learn California housing dataset and their mean square error. Neural Networks are designated NN and their architecture is described in parentheses with the first number being the input layer size and the final number the output layer size. The linear regression MSE was taken for the entire dataset while for the NN it was taken for the test dataset portion.

As can be seen from table ??, the linear regression performed significantly worse than the neural networks. As we know from estimation theory, the MSE of an estimator can be written as

$$MSE = \text{Var} + \text{Bias}^2 \quad (31)$$

which for linear regression is more strongly affected by the bias, as it is not strongly affected by small changes in the data, since it can't approximate non-linear relationships in the data. For neural networks, the bias tends to be lower as they can fit data with more complex relationships but is more perceptible to variance increasing the MSE, such as in the case of overfitting.

Looking at the model NN (8-32-64-16-1) in the jupyter notebook we can see that it achieved a final training loss of 0.208 but its test loss was 0.281. This difference of 35% is quite significant and indicates possible overfitting. Additionally, model NN (8-262144-1) exhibited unstable behavior with its loss fluctuating wildly. In comparison, the best models' test loss was similar to their training and evaluation losses.